

Week 2: Tabular Data

LSE MY472: Data for Data Scientists

<https://lse-my472.github.io/>

Autumn Term 2024

Ryan Hübert

Plan for today

- This week's git quick lesson: forking the website
- “Tidy data” and reshaping data in R
- Some good practises for code in (research) projects
- Coding

This week's `git` quick lesson: forking the website

Forking lse-my472.github.io

We discussed git branches and forks last week

You saw how to create a branch in your own repo

A **fork** is a copy of a repo that is its own repo

- Quite similar to a branch, but it is *not* within the same repo
- Pushed changes to a fork do not affect original remote repo
- Allows you to have a completely separate copy of a repo with your own changes while also pulling changes from the original remote as needed

Let's create a fork of the course website and clone locally

- Creates a copy of all the files in your own repo
- You can pull updates to your local copies every week

Forking lse-my472.github.io

1. Go to <https://github.com/lse-my472/lse-my472.github.io> and click Fork
2. Alter settings (optional) and click Create fork
3. Clone your fork to your computer using the usual process
4. Configure your computer so that it knows where the original remote repo is, and name it upstream:

```
git remote add upstream https://github.com/lse-my472/lse-my472.github.io.git
```

Important note: if you want to avoid headaches for now, do not make changes to the local copies of the files on your computer, as this may create merge conflicts

Syncing updates

Whenever you want to get updates (e.g., before lectures and seminars), sync your fork one of two ways

Option 1

1. On your computer, go to the command line, cd to your local copy of the fork
2. Fetch the changes from the original remote:
`git fetch upstream`
3. Check that you're on your master branch:
`git checkout master`
4. Merge the fetched changes from upstream into your fork: `git merge upstream/master`

Syncing updates

Whenever you want to get updates (e.g., before lectures and seminars), sync your fork one of two ways

Option 2

1. Go to your fork on GitHub and click “Sync Fork”
2. On your computer, go to the command line, `cd` to your local copy of the fork
3. Then, run `git pull` to bring changes from the GitHub copy of your fork to your local computer's copy

“Tidy data” and reshaping data in R

What is data?

- We can distinguish **data** ...
 - Representations, symbols, variables
 - E.g. “All”, “the”, “world’s”, “a”, “stage”
- ... from **information**
 - The meaning and context we gain from organizing and assembling data
 - E.g. “All the world’s a stage”
- Data are a convenient way of *storing* and *transmitting* information
 - See Caleb Sharf’s “The Ascent of Information” (2021)

Shapes of data

- One very common form of data is tabular
- Examples of other forms: Raw texts, key-value and array structures such as JSON files
- This week: How to organise and process tabular data (in R)
- Helpful to think about an ideal format of tabular data first

Comparing data structures

| | tweet | tags | lists | list_descriptions |
|---|---------------------------------------|------------|------------------------|-----------------------|
| 1 | I think MY472 is the best course ever | LSE, MY472 | asds_tweets, publicity | my thoughts, to share |

```
{
  "tweet" : {
    "message" : "I think MY472 is the best course ever",
    "tags" : ["LSE", "MY472"],
    "lists" : [
      {
        "name" : "asds_tweets",
        "description" : "my thoughts"
      },
      {
        "name" : "publicity",
        "description" : "to share"
      }
    ]
  }
}
```

“Tidy” data (Hadley Wickham)

Three rules:

1. Each variable must have its own column
2. Each observation must have its own row
3. Each value must have its own cell

| country | year | cases | population |
|-------------|------|--------|------------|
| Afghanistan | 1999 | 1845 | 15467071 |
| Afghanistan | 2000 | 2666 | 2055360 |
| Brazil | 1999 | 37737 | 172036362 |
| Brazil | 2000 | 89488 | 17404898 |
| China | 1999 | 213258 | 1272015272 |
| China | 2000 | 213666 | 128005583 |

variables

| country | year | cases | population |
|---------|------|-------|------------|
| ← | ← | ← | ← |
| ← | ← | ← | ← |
| ← | ← | ← | ← |
| ← | ← | ← | ← |
| ← | ← | ← | ← |
| ← | ← | ← | ← |

observations

| country | year | cases | population |
|---------|------|-------|------------|
| ○ | ○ | ○ | ○ |
| ○ | ○ | ○ | ○ |
| ○ | ○ | ○ | ○ |
| ○ | ○ | ○ | ○ |
| ○ | ○ | ○ | ○ |
| ○ | ○ | ○ | ○ |

values

Remainder of section based on

<https://r4ds.had.co.nz/tidy-data.html>

How will we work with “tidy” tabular data in R?

- In this course, we will use the `tidyverse` collection of packages to work with tabular data
- This package contains a bunch of useful tools; it's worth familiarising yourself at <https://www.tidyverse.org/>
- In the `tidyverse`, tabular data is stored in a `tibble`
- There are other ways to work with tabular data, e.g. “base R”
 - In base R, tabular data is stored as a `data.frame` or even `matrix`
- We don't care if you'd rather use base R for your own work, but we will mostly use `tidyverse` in class
 - We expect your turned in assignments to replicate whatever `tidyverse` would do

What does “tidy” data look like in R?

Here's a chunk of an example **dataset** (as a **tibble**)

```
table1
```

```
#> # A tibble: 6 × 4
```

```
#>   country      year  cases population
```

```
#>   <chr>      <dbl> <dbl>      <dbl>
```

```
#> 1 Afghanistan 1999     745  19987071
```

```
#> 2 Afghanistan 2000    2666  20595360
```

```
#> 3 Brazil      1999   37737  172006362
```

```
#> 4 Brazil      2000   80488  174504898
```

```
#> 5 China       1999  212258  1272915272
```

```
#> 6 China       2000  213766  1280428583
```

What can go wrong?

Datasets where columns represent values of a variable:

```
table4a
```

```
#> # A tibble: 3 x 3  
#>   country      `1999` `2000`  
#> * <chr>      <int>  <int>  
#> 1 Afghanistan    745   2666  
#> 2 Brazil         37737  80488  
#> 3 China          212258 213766
```

(There is another problem beyond the structure of the table...)

How to fix it?

We need to **pivot** those columns into a new pair of variables:

```
table4a %>%  
  pivot_longer(c(`1999`, `2000`), names_to = "year",  
    values_to = "cases")  
#> # A tibble: 6 x 3  
#>   country      year  cases  
#>   <chr>      <chr> <int>  
#> 1 Afghanistan 1999     745  
#> 2 Afghanistan 2000    2666  
#> 3 Brazil      1999   37737  
#> 4 Brazil      2000   80488  
#> 5 China       1999  212258  
#> 6 China       2000  213766
```


What is happening here?

We switched from **wide** to **long** format:

| country | year | cases | country | 1999 | 2000 |
|-------------|------|--------|-------------|--------|--------|
| Afghanistan | 1999 | 745 | Afghanistan | 745 | 2666 |
| Afghanistan | 2000 | 2666 | Brazil | 37737 | 80488 |
| Brazil | 1999 | 37737 | China | 212258 | 213766 |
| Brazil | 2000 | 80488 | | | |
| China | 1999 | 212258 | | | |
| China | 2000 | 213766 | | | |

table4

What else can go wrong?

Datasets where observations are scattered across multiple rows:

```
table2
```

```
#> # A tibble: 12 x 4
```

```
#>   country      year type      count
```

```
#>   <chr>      <int> <chr>    <int>
```

```
#> 1 Afghanistan  1999 cases      745
```

```
#> 2 Afghanistan  1999 population 19987071
```

```
#> 3 Afghanistan  2000 cases      2666
```

```
#> 4 Afghanistan  2000 population 20595360
```

```
#> 5 Brazil      1999 cases      37737
```

```
#> 6 Brazil      1999 population 172006362
```

```
#> # ... with 6 more rows
```

How to fix it?

We need to **pivot** those rows into a new pair of columns:

```
table2 %>%  
  pivot_wider(names_from = type, values_from = count)  
#> # A tibble: 6 x 4  
#>   country      year  cases population  
#>   <chr>      <int>  <int>      <int>  
#> 1 Afghanistan  1999     745   19987071  
#> 2 Afghanistan  2000    2666   20595360  
#> 3 Brazil       1999   37737   172006362  
#> 4 Brazil       2000   80488   174504898  
#> 5 China        1999  212258  1272915272  
#> 6 China        2000  213766  1280428583
```

What is happening here?

We switched from **long** to **wide** format:

| country | year | type | count |
|-------------|------|------------|------------|
| Afghanistan | 1999 | cases | 745 |
| Afghanistan | 1999 | population | 19987071 |
| Afghanistan | 2000 | cases | 2663 |
| Afghanistan | 2000 | population | 20595360 |
| Brazil | 1999 | cases | 37737 |
| Brazil | 1999 | population | 172006362 |
| Brazil | 2000 | cases | 80488 |
| Brazil | 2000 | population | 174504898 |
| China | 1999 | cases | 212258 |
| China | 1999 | population | 1272915272 |
| China | 2000 | cases | 213766 |
| China | 2000 | population | 1280428583 |

table2

| country | year | cases | population |
|-------------|------|--------|------------|
| Afghanistan | 1999 | 745 | 19987071 |
| Afghanistan | 2000 | 2663 | 20595360 |
| Brazil | 1999 | 37737 | 172006362 |
| Brazil | 2000 | 80488 | 174504898 |
| China | 1999 | 212258 | 1272915272 |
| China | 2000 | 213766 | 1280428583 |

Storing tabular data

Common file formats for storing tabular data:

- Comma-separated values (`.csv`) – ubiquitous and simple
 - Each *line* is an observation
 - Each variable value is separated by a comma
- Application specific (proprietary) formats (`.dta`, `.sav`, `.xls` etc.)
 - Can allow for richer representations including meta-data
 - More complex, and not necessarily human-readable
 - Can have explicit data limits (e.g. see Public Health England's *use of .xls spreadsheets*)

Often choice is dictated by the source (and size) of the data

- Packages like `haven` allow for reading in non-csv formats in R

Be very careful when storing data!

Some good practises for code in (research)
projects

Good practices in scientific computing

- Before we continue with examples of processing tabular data in R, it is helpful to spend some time early in this course with a brief discussion of good coding practises
- Based on Nagler (1995) “Coding Style and Good Computing Practices” (PS) and Wilson *et al* (2017) “Good Enough Practices in Scientific Computing” (PLOS Comput Biol)

Good practices in scientific computing

Why care?

→ Yourself

- Much lower chance of unnoticed bugs
- Future self will be grateful: “Yourself from 3 months ago doesn’t answer emails”
- More efficient research, avoid retracing own steps

→ Others

- Keep good records of what you did so that others can understand it
- **Replication** is a key part of science

Summary of some good practices

1. Safe and efficient data management
2. Well organised and documented code
3. Organised collaboration
4. One project = one repository
5. Track changes
6. Manuscripts as part of the analysis

1. Data management

- Save raw data as originally generated
- Create the data you would like to see, e.g.
 - If possible, open and non-proprietary formats such as .csv
 - Informative variable names instead of V322
 - Informative file names that contain metadata:
e.g. 05-alaska.csv instead of state5.csv
- Record all steps used to process data and store intermediate data files if computationally intensive (easier to rerun parts of a data analysis pipeline)
- Separate data manipulation from data analysis
- Prepare README with “codebook” of all variables
- Periodic backups (or Dropbox, Google Drive, etc.)
- Sanity checks: Summary statistics after data manipulation

2. Well organised and documented code

- Number scripts based on execution order
 - e.g. 01-clean-data.R, 02-recode-variables.R, 03-run-regression.R, 04-produce-figures.R...
- Write an explanatory note at the start of each script
 - Author, date of last update, purpose, inputs and outputs, other relevant notes
- Rules of thumb for modular code
 1. Any task you run more than once should be a function (with a meaningful name!)
 2. Many functions can be relatively short
 3. Can separate functions from execution (e.g. in functions.R file and then use `source(functions.R)` to load functions into current environment
- Try to keep it simple rather than too clever
- Add informative comments before blocks of code

3. Organised collaboration

- Create a README file with an overview of the project: Title, brief description, contact information, structure of folder
- Shared to-do list with tasks and deadlines
- Choose one person as corresponding author / point of contact / note taker
- Split code into multiple scripts to avoid simultaneous edits
- GitHub, Overleaf, Google Docs, etc. to collaborate in writing of manuscript

4. One project = one repository

Logical and consistent folder structure:

- `code` or `src` for all scripts
- `data` for raw data
- `temp` for temporary data files
- `output` or `results` for final data files and tables
- `figures` or `plots` for figures produced by scripts
- `manuscript` for text of paper
- `docs` for any additional documentation

5 & 6. Track changes; producing manuscript

- Ideally: Use version control (e.g. Git/GitHub)
- Manual approach: Keep dates versions of code & manuscript, and a `changelog` file with list of changes
- Dropbox also has some basic version control built-in
- Avoid typos and copy & paste errors
 - E.g., tables and figures can be produced in scripts and compiled directly into manuscript with Rmarkdown or \LaTeX

Examples

Barberá (2014) ([link to paper](#)):

- Replication materials (code and data)
- Code on GitHub

Thomas Leeper (2017) ([link to paper](#)):

- Replication materials (code and data)

Copus, Hübert and Pellaton (2024) ([link to paper](#)):

- Replication materials (code and data)

Also see John Myles White's [ProjectTemplate](#) R package

Coding

Coding

→ 01-conditionals-loops-functions.Rmd

→ 02-processing-data.Rmd